

Alexander Reinthal

📍 Gothenburg, Sweden ✉ email@reinthal.me ☎ 070-918 68 50 🌐 reinthal.me 📄 alexander-reinthal
🔗 reinthal

An independent AI safety researcher with a taste for work at the intersection of interpretability, misuse risks, and model organisms of misalignment. I bring 7 years of industry experience across cyber-security, data, and ML engineering.

Research Projects

[Aletheia's Quest: Lie detection Hackathon](#)

July 2026 – Aug 2026

Worked on the EleutherAI Team where we developed a hybrid method for lie detection that generalized well over 27 different datasets, 3 different model families and many different types of lies. My contributions were developing new datasets and validating data generation and labeling practices of previous research. I also wrote parts of the blog post which we will be releasing next week as part of the awards ceremony. We placed top 3.

[Can lying probes flag a Chinese model's deceptive answers?](#)

Jan 2026 – June 2026

Found that deception probes (Apollo, 2025) are unreliable at detecting questionable answers from Chinese models on politically sensitive questions.

- Extension of ARENA 7.0 weekend hackathon, 1st place. Supported by Bluedot Impact rapid grant.

[About Emergent Misalignment](#)

June 2026 – present

Found that emergent misalignment (EM) rate is negatively correlated with the capability benchmarks GSM8K (Pearson -0.98) and MMLU (Pearson -0.94).

- Extension of ARENA 7.0 capstone, 3rd place.

[Changing the prompt for instrumental convergence evals](#)

May 2026 – June 2026

Removing pressure phrases increases misalignment rate for capability prompts ([He et al. 2025](#)) but decreases it for propensity prompts ([Hopman et al. 2026](#)) for Gemini 3.1 Pro Preview.

[Detecting Piecewise Cyber Espionage in Model APIs](#)

Nov 2025 – Nov 2025

Proof of concept for detecting malicious activity across multiple agent sessions.

- Research continued through SPAR by David Williams-King (ERA) and Linh Le (Mila). Apart Research sprint, 4th place of 641 submissions.

Experience

AI Safety Researcher

May 2026 – present

Independent

- My research focuses on AI control and scalable oversight methods to be applied throughout the model lifecycle. Supported by Coefficient Giving's Career Transition grant for reducing global catastrophic risks.

Research mentor

Mar 2026 – present

Bluedot Impact

- Facilitator and research mentor for the Technical AI Safety course and project.

Alum

Jan 2026 – Feb 2026

ARENA 7.0

- Five-week intensive curriculum on transformer internals, mechanistic interpretability, RLHF, evals, JAX, and the Inspect framework.

Forward-deployed Engineer / DevSecOps / Tech Lead / Data Scientist
Knowit Solutions Cocreate

Gothenburg, Sweden
Feb 2022 – present

- Lead secure deployment of data platforms across multiple clients, some under protective security contracts ([Säkerhetsskyddslagen 2018:585](#)).
- Pioneered LLM evals at the company. Implemented an OAuth-authenticated MCP gateway to secure agentic workflows.

Security Analyst
NTT Security

Gothenburg, Sweden
Apr 2019 – Jan 2022

- Deployed an ML classifier that reduced analysis time of malicious domains from 2 hours to 10 minutes per analyst per shift.

Python Software Engineer
Ericsson

Gothenburg, Sweden
June 2017 – Oct 2018

- First job where I led a team building a log-parsing and analytics tool that got funded and grew into a dedicated engineering team.

Publications

Data Modeling for Predicting Software Exploits

Nov 2018

Peer-reviewed paper predicting which CVEs become exploited in the wild. Adopted by Recorded Future's Vulnerability Intelligence product.

Reinthal, A., Filippakis, E., Almgren, M.. [10.1007/978-3-030-03638-6_21](#) (Springer LNCS: Nordic Conference on Secure IT Systems)

Open Source and Personal Projects

Casually Jailbreaking Gemini 2.5 Flash

2026 – 2026

Jailbroke Gemini 2.5 Flash with a few-shot in-context continuation attack, eliciting harmful recommendations.

Open Source Contributions

github.com/reinthal

- [Inspect AI PR #3987](#) Fixed Gemini 3.1 Pro + OpenRouter + Reasoning issue
- [Dagster PR #24188](#) (ML pipeline orchestration) and [DLT Hub PR #594](#) (ingestion reliability)

Skills

Machine Learning & AI: PyTorch, JAX (learned during ARENA 7.0), Einops, XGBoost, SciKit Learn, Jupyter; LLM evals & observability (Inspect, LangChain, Langfuse, Arize Phoenix); linear probes, SFT, automated red-teaming

Programming Languages: Python, JavaScript, SQL, Rust, C/C++, R, Nix, Terraform

Platforms & Tools: Modal, Runpod, Openweights, Weights & Biases, Spark, Docker, Kubernetes, AWS, Azure

Certifications: Bluedot Impact — Technical AI Safety (2026-03); AGI Strategy (2025-10)

Education

Chalmers University of Technology
MS in Engineering Physics

Gothenburg, Sweden
Sept 2016 – June 2018

- Master thesis got published and put into production at Recorded Future.

- Switched from computer science undergraduate to engineering physics graduate program and performed better.

Gothenburg University

BS in Computer Science

Gothenburg, Sweden

Sept 2013 – June 2016

- Thesis on finding common dense subgraphs using linear programming has 4 citations.