

Alexander Reinthal

📍 Gothenburg, Sweden ✉ email@reinthal.me ☎ 070-918 68 50 🌐 reinthal.me 📄 alexander-reinthal
🔗 reinthal

An independent AI safety researcher with a taste for work at the intersection of interpretability, misuse risks, and model organisms of misalignment. I bring 7 years of industry experience across cyber-security, data, and ML engineering.

Research Projects

[Aletheia's Quest: Lie detection Hackathon](#)

July 2026 – Aug 2026

Competed on the EleutherAI Team throughout the competition and ranked no. 1 for both black-box and white-box detection methods. Blog post coming soon.

[Can lying probes flag a Chinese model's deceptive answers?](#)

Jan 2026 – June 2026

Found that deception probes (Apollo, 2025) are unreliable at detecting questionable answers from Chinese models on politically sensitive questions.

- Extension of ARENA 7.0 weekend hackathon, 1st place. Supported by Bluedot Impact rapid grant.

[About Emergent Misalignment](#)

June 2026 – present

Found that emergent misalignment (EM) rate is negatively correlated with the capability benchmarks GSM8K (Pearson -0.98) and MMLU (Pearson -0.94).

- Extension of ARENA 7.0 capstone, 3rd place.

[Changing the prompt for instrumental convergence evals](#)

May 2026 – June 2026

Removing pressure phrases increases misalignment rate for capability prompts ([He et al. 2025](#)) but decreases it for propensity prompts ([Hopman et al. 2026](#)) for Gemini 3.1 Pro Preview.

[Detecting Piecewise Cyber Espionage in Model APIs](#)

Nov 2025 – Nov 2025

Proof of concept for detecting malicious activity across multiple agent sessions.

- Research continued through SPAR by David Williams-King (ERA) and Linh Le (Mila). Apart Research sprint, 4th place of 641 submissions.

Experience

AI Safety Researcher

May 2026 – present

Independent

- Supported by Coefficient Giving's Career Transition grant for reducing global catastrophic risks.

Research mentor

Mar 2026 – present

Bluedot Impact

- Facilitator and research mentor for the Technical AI Safety course and project.

Alum

Jan 2026 – Feb 2026

ARENA 7.0

- Five-week intensive curriculum on transformer internals, mechanistic interpretability, RLHF, evals, JAX, and the Inspect framework.

Forward-deployed Engineer / DevSecOps / Tech Lead / Data Scientist

Knowit Solutions Cocreate

Gothenburg, Sweden

Feb 2022 – present

- Lead secure deployment of data platforms across multiple clients, some under protective security contracts ([Säkerhetsskyddslagen 2018:585](#)).
- Pioneered LLM evals at the company. Implemented an OAuth-authenticated MCP gateway to secure agentic workflows.

Security Analyst

NTT Security

Gothenburg, Sweden

Apr 2019 – Jan 2022

- Deployed an ML classifier that reduced analysis time of malicious domains from 2 hours to 10 minutes per analyst per shift.

Python Software Engineer

Ericsson

Gothenburg, Sweden

June 2017 – Oct 2018

- First job where I led a team building a log-parsing and analytics tool that got funded and grew into a dedicated engineering team.

Publications

Data Modeling for Predicting Software Exploits

Nov 2018

Peer-reviewed paper predicting which CVEs become exploited in the wild. Adopted by Recorded Future's Vulnerability Intelligence product.

Reinthal, A., Filippakis, E., Almgren, M.. [10.1007/978-3-030-03638-6_21](#) (Springer LNCS: Nordic Conference on Secure IT Systems)

Open Source and Personal Projects

[Casually Jailbreaking Gemini 2.5 Flash](#)

2026 – 2026

Jailbroke Gemini 2.5 Flash with a few-shot in-context continuation attack, eliciting harmful recommendations.

Open Source Contributions

github.com/reinthal

- [Inspect AI PR #3987](#) Fixed Gemini 3.1 Pro + OpenRouter + Reasoning issue
- [Dagster PR #24188](#) (ML pipeline orchestration) and [DLT Hub PR #594](#) (ingestion reliability)

Skills

Machine Learning & AI: PyTorch, JAX (learned during ARENA 7.0), Einops, XGBoost, SciKit Learn, Jupyter; LLM evals & observability (Inspect, LangChain, Langfuse, Arize Phoenix); linear probes, SFT, automated red-teaming

Programming Languages: Python, JavaScript, SQL, Rust, C/C++, R, Nix, Terraform

Platforms & Tools: Modal, Runpod, Openweights, Weights & Biases, Spark, Docker, Kubernetes, AWS, Azure

Certifications: Bluedot Impact — Technical AI Safety (2026-03); AGI Strategy (2025-10)

Education

Chalmers University of Technology

MS in Engineering Physics

Gothenburg, Sweden

Sept 2016 – June 2018

- Master thesis got published and put into production at Recorded Future.
- Switched from computer science undergraduate to engineering physics graduate program and performed better.

- Thesis on finding common dense subgraphs using linear programming has 4 citations.