

Alexander Reinthal

📍 Gothenburg, Sweden ✉ email@reinthal.me ☎ 073-036 28 78 🔗 reinthal.me in alexander-reinthal
 🌐 reinthal

Alexander Reinthal

Senior Data Engineer, Independent Contributor & Open-Source Advocate with expertise in real-time data ingestion, data modeling, and orchestrating cloud-based data workloads. Strong background in **Apache Spark, Flink & Iceberg, Databricks & Unity Catalog, and Snowflake**, similar to **AWS Glue, S3, Redshift and EMR-equivalent architectures**. Proven open-source contributions to Dagster & DLT Hub. Passionate about developer experience, mentoring engineers in distributed systems, and ML-driven analytics.

Experience

Knowit Solutions Cocreate, Data Engineer / Tech Lead / Delivery Manager

Gothenburg, Sweden
Feb 2022 – present

- Worked remotely leading development of an asset-centric data modeling approach for Databricks & Delta Live Tables, enhancing streaming data reliability and developer experience.
- Led an agile R&D team to build real-time data ingestion pipelines handling massive volumes of CI-pipeline logs using **Apache Flink (AWS Kinesis/Data Streams equivalent) & Spark Structured Streaming**.
- Designed scalable data pipelines on **Databricks & Snowflake, similar to AWS EMR & Amazon Redshift**, enabling advanced analytics workloads.
- Served as the technical lead, engaging with a client to design and implement a modern data platform leveraging Snowflake, Dagster, and dbt.
- Mentored junior engineers in Python, SQL, DAG-based orchestration, and modern data engineering best practices. Led workshops on aligning technical solutions with business objectives.
- Collaborated with stakeholders to refine data governance, mediate upstream data contracts, and improve data visibility using effective metadata management.

NTT Security, Security Analyst

Gothenburg, Sweden
Apr 2019 – Jan 2022

- Developed predictive analytics tooling for detecting malicious network traffic, leveraging Python, machine learning, and distributed processing (precursor to modern real-time data pipelines).
- Automated security alert processing, optimizing real-time event analysis and reducing response times.
- Designed Python-based data ingestion and processing workflows, akin to real-time ETL pipelines in modern data engineering.
- Provided mentorship and security training, educating analysts on threat intelligence, anomaly detection, and alert triage automation.

Ericsson, Python Software Engineer

Gothenburg, Sweden
June 2017 – Oct 2018

- Led a team to develop a log parsing and analytics tool that scaled into a dedicated engineering team at Ericsson.
- Automated data workflow optimizations in Python, improving efficiency in log-based machine-generated data.

Projects

HowIgotPwned

GitHub [↗](#)

- Developed a modern data lakehouse solution leveraging Apache Iceberg, Flink,

Spark, Nessie (similar to AWS Glue), MinIO (AWS S3 equivalent), and Dagster to analyze high-volume data.

- Prototyped real-time ingestion pipelines using Apache Flink and Iceberg tables stored on AWS S3, demonstrating cloud-scale flexibility.

Datalake Stack (On-Prem Equivalent of AWS Data Architecture)

[GitHub](#)

- Designed and implemented an on-premise data lake solution for Kubernetes using Dagster, Nessie, Apache Iceberg, MinIO, Apache Flink and Spark. Similar to AWS Glue, AWS S3 Tables and AWS S3.
- Inspired work to scale our open source data platform to multiple customers.

Open Source Contributions

[GitHub](#)

- Dagster Contribution: [PR #24188](#) – Improved workflow orchestration, scheduling, and monitoring of data pipelines.
- DLT Hub Contribution: [PR #594](#) – Mentored junior engineers, contributing Python data ingestion enhancements.

Technologies

Data Architectures: Data Modeling, Data Warehousing, Data Lakehouse Solutions

Semantic Layers: dbt Cloud, Cube.dev

Data Catalogs: *Polaris, Snowflake, Unity & Nessie (Similar to Amazon Glue)*

Big Data & Streaming: *Apache Flink, Apache Spark (AWS EMR equivalent)*, RabbitMQ

Data Engineering Tools & Platforms: *Databricks & Delta Live Tables, Snowflake (similar to Amazon Athena)*, dbt, Dagster

Languages: Python, SQL, Java, Terraform, Nix, Rust, C++, Assembly, R

Cloud & Platforms: MinIO & Azure Blob Storage (AWS S3 equivalent), Azure, Kubernetes, Proxmox

Infrastructure & DevOps: Docker, Podman, FluxCD, Git, CI/CD pipelines, Terraform, Helm

Business Intelligence Tools: PowerBI, Apache Superset, Lightdash

Education

MS **Chalmers Technical University**, Engineering Physics Sept 2016 – June 2018

- Coursework: Statistical Physics, Neural Networks and Machine Learning

BS **Göteborg University**, Computer Science Sept 2013 – June 2016

- Coursework: Algorithms, Testing, Debugging and Verification, Theoretical Computer Science, Operating Systems, Cryptography, Cyber Security

Publications

Data Modelling for Predicting Exploits

Nov 2018

Reinthal, A, Filippakis, E., Almgren, M.

10.1007/978 – 3 – 030 – 03638 – 6₂1 [↗](#)